

Contents

| | |
|---|---|
| Notes on Hands-on Workshop of ingest nodes | 1 |
| Agenda | 1 |
| Definitions | 2 |
| Quick setup instructions for OSX (Mac) | 2 |
| Cookiecutter run | 2 |
| Update your <code>.env</code> file | 3 |
| Copy an existing DBT project to use for now | 3 |
| Adjust the project and model name to match your <code>dummy_jaffle_</code> name in the <code>dbt_project.yml</code> file | 3 |
| Source aliases and log into ACR, pull images | 3 |
| Check that you are able to connect to the Databricks SQL Warehouse | 3 |
| Check that <code>dbt debug</code> succeeds | 4 |
| Optionally change the log file format to TEXT since we are not pushing logs already and this is more readable | 4 |
| From here you can try out any of the modalities: | 4 |
| Feedback questions | 4 |
| How do we specify the source and sink tables for a given transfor- mation? | 4 |
| AR: What are some of the best resources for us to learn about DBT? | 5 |
| AR: Can we see a simple example of a simple data transformation from end to end? | 5 |
| AR: Can we expand on the simple example, by chaining two simple transformations together using pipelines. | 5 |
| Observability | 5 |

Notes on Hands-on Workshop of ingest nodes

These notes should be read in conjunction with the session recording.

Agenda

- Review of last session and context setting for this one
 - Previous session was focussed on the “Why”, today we will focus on the “How”.
- Build your own standard DBT project: DBT-based ingest node “Dummy Jaffle”.
- Addressing the feedback questions
- Demonstrating how to share data with Unity Catalog
- Demo of DBT project with `automate_dv` DBT macro extension: “Dummy Bridgetown”.

We will demonstrate aspects of:

- DBT-based ingest node
 - best suited for OLAP use case

- can ingest data through `dbt seed`
- can ingest data from cloud storage through `COPY INTO` statement
- can model data using DBT models
- can use macro packages such as `automate_dv` to simplify data vault building

Definitions

Quick setup instructions for OSX (Mac)

To get set up: * Docker Desktop. Choose Apple Chip/Intel depending on your Mac and install. * `brew install git` if you haven't yet * `xcode-select --install` if you haven't installed Xcode command line tools yet (Might take a while to download/install) * Install cookiecutter: `run export PATH=$HOME/.local/bin:$PATH` and then `brew install cookiecutter` * For pulling docker images from Azure Container Registry we need to authenticate via the Azure CLI, so please install: `brew update && brew install azure-cli` * Generate Databricks token, in User Settings -> Access Tokens

Cookiecutter run

- All of the answers below to the Cookiecutter questions are obtained from ADO and Databricks.
- The ADO config file is something you have to setup in your home directory.
- The ADO config file is only required for the ADO Pipeline automation, at the time when you want to setup the build, release and schedule pipelines.

```
cookiecutter git@ssh.dev.azure.com:v3/exploreai/CORE.Utilities/CORE.Meshnodes.BaseTemplates/
ingest_node_name [Ingest Node Template]: dummy_jaffle_kr
az_devops_repo [https://dev.azure.com/exploreai/CORE.Utilities/_git/CORE-POC1]: https://dev
Select data_substrate:
1 - databricks
2 - postgres
Choose from 1, 2 [1]: 1
databricks_workspace [https://adb-891777510264692.12.azuredatabricks.net/]: https://adb-4472
database_host [adb-891777510264692.12.azuredatabricks.net]: adb-4472170994427587.7.azuredat
databricks_sql_warehouse_path [/sql/1.0/endpoints/a077556ed384ed67]: /sql/1.0/warehouses/22
databricks_catalog [null]: dummy_jaffle
database_user [psqladmin]:
database_port [5432]: 443
database_threads [4]: 5
profile_name [data_vault]:
azure_storage_container_url [wasbs://adf-pipeline-demo@datavalidation.blob.core.windows.net,
ado_pipelines_build_agent_pool_name [CORE Pipelines]: Default
ado_pipelines_folder_name [CORE]: dummy_jaffle
ado_profile_name [CORE]: DEFAULT
scheduled_release_cron [0 6 * * *]:
```

```
-----  
Template clone successful  
-----
```

```
--> Attempting to move pipeline definitions to relevant location...  
    --> Moving files...  
    --> File move successful.  
--> Searching for ADO config file in home dir...  
? ADO config file detected. Would you like to alter/extend it? No  
User declined to modify ADO config file at /home/kerneels/.adocfg. Skipping file generation  
  
--> Generating `.env` file from supplied definition file...  
    --> Generation successful.  
    --> Please inspect/edit the generated file at `dummy_jaffle/.env` to ensure that th
```

Update your .env file

- Set your PAT for ACCESS_TOKEN.
- Set a SAS token as AZURE_SAS_TOKEN if you intend to read from blob storage. Ensure the token has read and list privileges.

Copy an existing DBT project to use for now

```
cp -fr dummy_jaffle/src/ dummy_jaffle_<initials>/
```

Adjust the project and model name to match your dummy_jaffle__ name in the dbt_project.yml file

- Find the dbt_project.yml file in dummy_jaffle_<initials>/src/ folder.
- Edit it and rename dummy_jaffle to dummy_jaffle_kr to match your new project.

Source aliases and log into ACR, pull images

```
cd dummy_jaffle_<initials>/  
source .bash_aliases  
ingest-acr-login  
ingest-pull  
ingest-build
```

Check that you are able to connect to the Databricks SQL Warehouse

```
ingest-dbsqlcli
```

Check that dbt debug succeeds

ingest-debug

Optionally change the log file format to TEXT since we are not pushing logs already and this is more readable

This is just editing the .env file in place with sed:

```
sed -i 's/DBT_LOG_FORMAT=JSON/DBT_LOG_FORMAT=TEXT/g' .env
```

From here you can try out any of the modalities:

- **ingest-seed** - to load any seed data via **dbt seed**
- **ingest-debug** - see if everything is setup properly via **dbt debug**
- **ingest-dbsqlcli** - connect to the Databricks Serverless SQL Warehouse or Standard Cluster where you can query using SQL (exit using the keyword **exit**)
- **ingest-exec** - run the entire DBT project (all models) via **dbt run**
- **ingest-test** - run all the tests via **dbt test**
- **ingest-docs** - generate documentation and serve it on <http://localhost:8080/>
- **ingest-cli** - enter into the running docker container from where you can issue ANY DBT command
- **ingest-other** - specify variable **MODE_SWITCH_OTHER** with ANY DBT command you would like to issue

Feedback questions

How do we specify the source and sink tables for a given transformation?

- The smallest unit of a transformation is a DBT model (we will just call it model onwards).
- A model is essentially something that takes data from one or more inputs and produce one output.
- The input data for a model is one of these:
 - any existing data already ingested/transformed perhaps by other systems - we call these sources and they are referenced via the `{{ source() }}` function.
 - * note that we have to define them in YML files so DBT knows about them
 - any data ingested via seeds referenced via the `{{ ref('model_name') }}` function.
 - any other model in the DAG referenced via the `{{ ref('model_name') }}` function.
- The output of a model is determined by the type of materialisation configured, and it can be a table or a view.

- We do not explicitly “save” into a table, but let DBT decide when it should persist or not

AR: What are some of the best resources for us to learn about DBT?

- Given these already but will list them no problem!

AR: Can we see a simple example of a simple data transformation from end to end?

AR: Can we expand on the simple example, by chaining two simple transformations together using pipelines.

Observability

- We can emit logs to Azure Log Analytics Workspace (LAW).
- Ensure all the LAW-related variables are set.
- If some of the LAW-related variables are not set, the system will inform you of this and no logs can be sent.
- The log data goes here.
- Take a look at this dashboard built off of the data.
- you can alter it and query the log data using KQL.

These notes should be read in conjunction with the session recording.